

# NAÏVE BAYES


---

# Generative vs. Discriminative Classifiers

Training classifiers involves estimating  $f: X \rightarrow Y$ , or  $P(Y|X)$

Discriminative classifiers (also called ‘informative’ by Rubinstein&Hastie):

1. Assume some functional form for  $P(Y|X)$
2. Estimate parameters of  $P(Y|X)$  directly from training data




# Generative vs. Discriminative Classifiers

Training classifiers involves estimating  $f: X \rightarrow Y$ , or  $P(Y|X)$

Generative classifiers

1. Assume some functional form for  $P(X|Y)$ ,  $P(X)$
2. Estimate parameters of  $P(X|Y)$ ,  $P(X)$  directly from training data
3. Model predictions based on  $P(Y|X = x_i)$



# Maximum A Posteriori (MAP) Classifier

- Given the data feature vector  $\mathbf{x}$ , we would like to find the class with the largest probability:

$$\hat{C} \triangleq \arg \max_C p(C|\mathbf{x})$$

- To accomplish this, we iterate through all possible classes  $C_1, C_2, \dots, C_K$  and evaluate the quantity  $p(C_i|\mathbf{x})$ , and pick the class that has the largest probability.

# MAP Classifier: Bayes Rule

- How do we compute  $p(C_i|\mathbf{x})$  ?      Apply Bayes' Rule!

$$p(A|B)p(B) = p(B|A)p(A) \Rightarrow p(A|B) = \frac{p(B|A)p(A)}{p(B)}$$

likelihood      prior  
evidence

- Applying Bayes' rule to  $p(C_i | \mathbf{x})$  , we obtain:

$$p(C_i|\mathbf{x}) = \frac{p(\mathbf{x}|C_i)p(C_i)}{p(\mathbf{x})}$$

# MAP Classifier: Bayes Rule (2)

$$p(C_i|\mathbf{x}) = \frac{p(\mathbf{x}|C_i)p(C_i)}{p(\mathbf{x})}$$

- $p(C_i)$  is called the prior probability of a class
- $p(\mathbf{x}|C_i)$  is called the likelihood of the data (what is the probability of observing  $\mathbf{x}$  if the class was  $C_i$ )
- $p(\mathbf{x})$  is the probability of seeing the data so its simply a normalizing factor (applies to all  $C_i$ ), so doesn't affect which  $C_i$  attains MAP.

# MAP Classifier: Bayes Rule (3)

- To compute the normalizing factor, we use:

$$p(\mathbf{x}) = \sum_{i=1}^K p(\mathbf{x}|C_i)p(C_i)$$

So that

$$\sum_{i=1}^K p(C_i|\mathbf{x}) = 1$$

Hence,

$$p(C_i|\mathbf{x}) = \frac{p(\mathbf{x}|C_i)p(C_i)}{\sum_{\ell=1}^K p(\mathbf{x}|C_{\ell})p(C_{\ell})}$$

But really, we don't need to calculate the denominator  $p(\mathbf{x})$ , so we can simply write it as

$$p(C_i|x) = p(x|C_i)p(C_i)$$

# MAP Classifier: Bayes Rule (4)

- Usually, the likelihood  $p(\mathbf{x}|C_i)$  is difficult to compute because it is N-dimensional (length of feature vector).
- This is because the distribution considers correlations between the features when computing the likelihood.
- We can write  $p(\mathbf{x}|C_i)$  equivalently as:

$$p(\mathbf{x}|C_i) = p(x_1, x_2, \dots, x_N|C_i)$$

which makes the dependence on individual features explicit.



# Naïve Bayes Classifier: Independence Assumption

- The Naïve Bayes classifier introduces one major assumption regarding the features: **independence**
- That is, the Naïve Bayes classifier assumes:

$$p(\mathbf{x}|C_i) = p(x_1, x_2, \dots, x_N|C_i) = \prod_{n=1}^N p(x_n|C_i)$$

- That is, the complicated likelihood  $p(\mathbf{x}|C_i)$  can now be factored into a product of N 1-dimensional likelihoods, which are easy to compute.
- It is good to note that independence implies that the features are not correlated (but generally not the other way around---so independence is a stronger assumption).

# Naïve Bayes Classifier

- With the independence assumption, the MAP classifier (now called the Naïve Bayes Classifier) can be written as:

$$p(C_i|\mathbf{x}) = \frac{\prod_{n=1}^N p(x_n|C_i)p(C_i)}{\sum_{\ell=1}^K \prod_{n=1}^N p(x_n|C_{\ell})p(C_{\ell})}$$

- It turns out that the Naïve Bayes classifier works very well on empirical datasets even if the independence assumption doesn't actually hold.

# Naïve Bayes Classifier

$$p(C_i|\mathbf{x}) = \prod_{n=1}^N p(x_n|C_i)p(C_i)$$

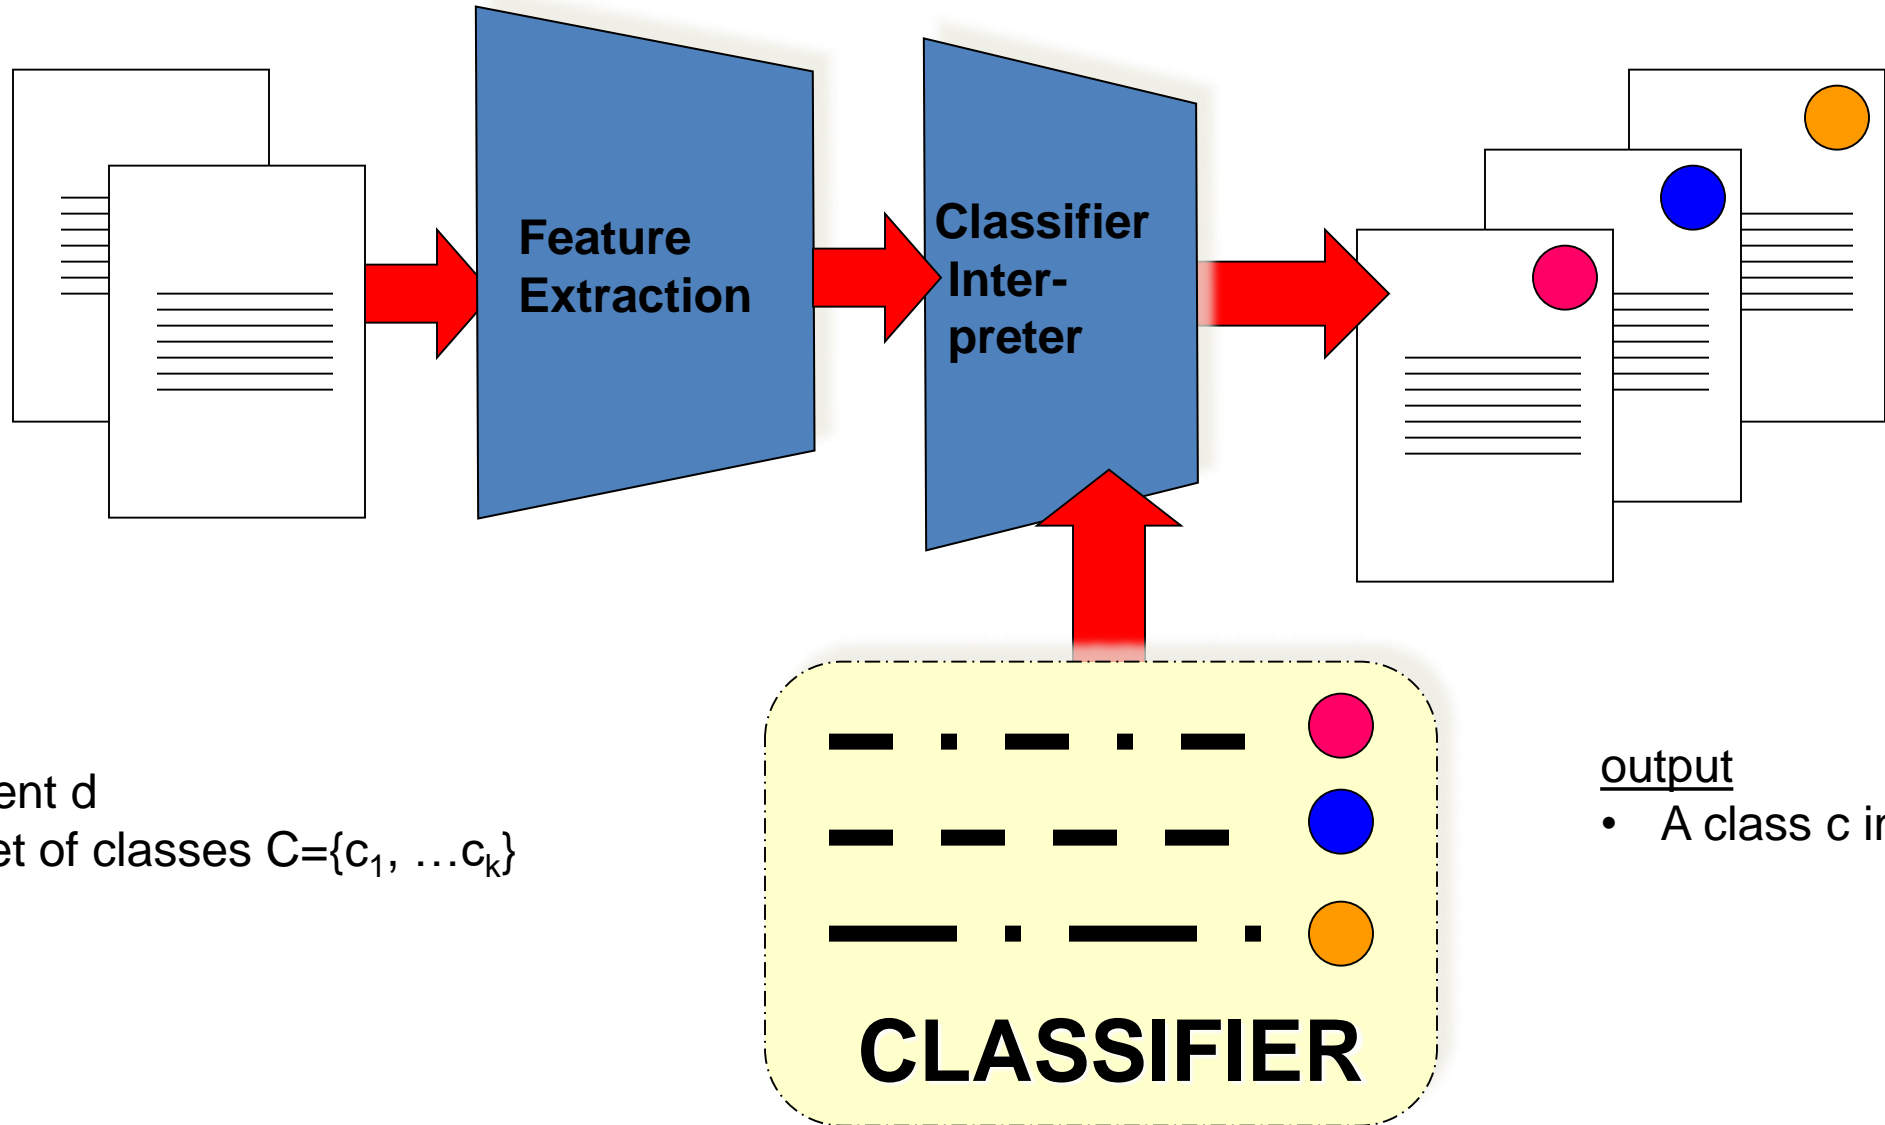
- What happens our training example does not have an instance that is present in the test data?
- To eliminate zeros, we use add-one or Laplace smoothing, which simply adds one to each count.

Now, let us see Spark documentation: <http://spark.apache.org/docs/latest/mllib-naive-bayes.html>

# EXAMPLES OF TEXT CATEGORIZATION

- Spam Detection
  - “spam” / “not spam”
- Assigning subject categories, topics, or genres
  - “finance” / “sports” / “asia”
- Sentiment analysis
  - “like” / “hate” / “neutral”
- Authorship identification
  - “Shakespeare” / “Marlowe” / “Ben Jonson”
  - The Federalist papers

# TEXT CATEGORIZATION



## Input

- A document  $d$
- A fixed set of classes  $C = \{c_1, \dots, c_k\}$

## output

- A class  $c$  in  $C$

# Text Classification Algorithms: Learning

- From training corpus, extract Vocabulary
- Calculate required  $P(c_j)$  and  $P(x_k | c_j)$  terms
  - For each  $c_j$  in  $C$  do
    - $docs_j \leftarrow$  subset of documents for which the target class is  $c_j$

$$P(c_j) \leftarrow \frac{|docs_j|}{|\text{total \# documents}|}$$

- $Text_j \leftarrow$  single document containing all  $docs_j$ 
  - for each word  $x_k$  in *Vocabulary*
    - $n_k \leftarrow$  number of occurrences of  $x_k$  in  $Text_j$

$$P(x_k | c_j) \leftarrow \frac{n_k + 1}{n + |Vocabulary|}$$

This is called the Bag of Words model.... But how do we select this vocabulary??

Can't possibly use all words!

Chi Square, TF/IDF, etc.

# Text Classification Algorithms: Classifying

- $positions \leftarrow$  all word positions in current document which contain tokens found in Vocabulary
- Return  $c_{NB}$ , where

$$c_{NB} = \operatorname{argmax}_{c_j \in C} P(c_j) \prod_{i \in positions} P(x_i | c_j)$$

# Naïve Bayes Posterior Probabilities

- Classification results of naïve Bayes (the class with maximum posterior probability) are usually fairly accurate.
- However, due to the inadequacy of the conditional independence assumption, the actual posterior-probability numerical estimates are not.
  - Output probabilities are generally very close to 0 or 1.